

# 10 Creating Functions

---

## ✓ Differentiate between a procedure and a function

Function difference: [8.5.1](#) - [8.5.2](#)

A function has the same structure as a procedure, except that:

- A function heading must include a **RETURN** clause, which specifies the data type of the value that the function returns. (A procedure heading cannot have a **RETURN** clause.)
- In the executable part of a function, every execution path must lead to a **RETURN** statement. Otherwise, the PL/SQL compiler issues a compile-time warning. (In a procedure, the **RETURN** statement is optional and not recommended)
- A function declaration can include these options:
  - **DETERMINISTIC**: Helps the optimizer avoid redundant function invocations
  - **PARALLEL\_ENABLE**: Enables the function for parallel execution, making it safe for use in concurrent sessions of parallel DML evaluations
  - **PIPELINED**: Makes a table function pipelined, for use as a row source
  - **RESULT\_CACHE**: Stores function results in the PL/SQL function result cache

### RETURN statement

- The **RETURN** statement immediately ends the execution of the subprogram or anonymous block that contains it
- A subprogram or anonymous block can contain multiple **RETURN** statements
- Note: The **RETURN** statement differs from the **RETURN** clause in a function heading, which specifies the data type of the return value
- syntax: **RETURN** [**expression**]
  - expression:
    - Optional when the **RETURN** statement is in a pipelined table function
    - Required when the **RETURN** statement is in any other function
    - Not allowed when the **RETURN** statement is in a procedure or anonymous block
    - The **RETURN** statement assigns the value of expression to the function identifier. Therefore, the data type of expression must be compatible with the data type in the **RETURN** clause of the function
- **RETURN statement in function**
  - In a function, every execution path must lead to a **RETURN** statement and every **RETURN** statement must specify an expression
  - The **RETURN** statement assigns the value of the expression to the function identifier and returns control to the invoker, where execution resumes immediately after the invocation
  - Note: In a pipelined table function, a **RETURN** statement need not specify an expression

- Every execution path in a function must lead to a **RETURN** statement (otherwise, the PL/SQL compiler issues compile-time warning **PLW-05005**)
- If a function has an execution path without **RETURN** (e.g. if statement with no else clause with return), then the function compiles with warning **PLW-05005: subprogram ... returns without value at line ...**

- **RETURN statement in procedure**

- In a procedure, the **RETURN** statement returns control to the invoker, where execution resumes immediately after the invocation
- The **RETURN** statement cannot specify an expression

- **RETURN statement in anonymous block**

- In an anonymous block, the **RETURN** statement exits its own block and all enclosing blocks
- The **RETURN** statement cannot specify an expression
- Example:

```
BEGIN
  BEGIN
    DBMS_OUTPUT.PUT_LINE('Inside inner block.');
```

**RETURN;**

```
    DBMS_OUTPUT.PUT_LINE('Unreachable statement.');
```

**END;**

```
  DBMS_OUTPUT.PUT_LINE('Inside outer block. Unreachable statement.');
```

**END;**

**/**

-- result:  
Inside inner block.

## ✓ Describe the uses of functions

[Function result cache: 8.12](#)

[Functions that SQL can invoke: 8.13](#)

### Function result cache

- When a PL/SQL function has the **RESULT\_CACHE** option, its results are cached in the shared global area (SGA) so sessions connected to the same instance can reuse these results when available
- Oracle Database automatically detects all data sources (tables and views) that are queried while a result-cached function is running. If changes to any of these data sources are committed, the cached result becomes invalid across all instances
- The best candidates for result-caching are functions that are invoked frequently but depend on information that changes infrequently or never
- Result-caching avoids redundant computations in recursive functions

## Enabling result-caching for a function

- To make a function result-cached, include the **RESULT\_CACHE** clause in the function declaration and definition
- The **RESULT\_CACHE** clause can appear in the following SQL statements: **CREATE FUNCTION**, **CREATE TYPE BODY**
- If you declare the function before defining it, you must also include the **RESULT\_CACHE** option in the function declaration
- syntax: **RESULT\_CACHE [ RELIES\_ON ( [data\_source [,...]] ) ]**
  - **RELIES\_ON**:
    - Specifies the data sources on which the results of the function depend. Each data\_source is the name of either a database table or view
    - This clause is deprecated. As of Oracle Database 12c, the database detects all data sources that are queried while a result-cached function is running, and **RELIES\_ON** clause does nothing
- restrictions:
  - **RESULT\_CACHE** is disallowed on functions with **OUT** or **IN OUT** parameters
  - **RESULT\_CACHE** is disallowed on functions with **IN** or **RETURN** parameter of (or containing) these types: **BLOB**, **CLOB**, **NCLOB**, **REF CURSOR**, Collection, Object, Record
  - **RESULT\_CACHE** is disallowed on function in an anonymous block
  - **RESULT\_CACHE** is disallowed on pipelined table function and nested function

## Developing applications with result\_cached functions

- When developing an application that uses a result-cached function, make no assumptions about the number of times the body of the function will run for a given set of parameter values
- Some situations in which the body of a result-cached function runs are:
  - The first time a session on this database instance invokes the function with these parameter values
  - When the cached result for these parameter values is invalid When a change to any data source on which the function depends is committed, the cached result becomes invalid
  - When the cached results for these parameter values have aged out If the system needs memory, it might discard the oldest cached values
  - When the function bypasses the cache

## Requirements for result-cached functions

- A result-cached PL/SQL function is safe if it always produces the same output for any input that it would produce were it not marked with **RESULT\_CACHE**
- This safety is only guaranteed if these conditions are met:
  - When the function is executed, it has no side effects
  - All tables that the function accesses are ordinary, non-SYS-owned permanent tables in the same database as the function
  - The function's result must be determined only by the vector of input actuals together with the committed content, at the current SCN, of the tables that it references
- It is recommended that a result-cached function also meet these criteria:
  - It does not depend on session-specific settings
  - It does not depend on session-specific application contexts

## Rules for a cache hit

- Each time a result-cached function is invoked with different parameter values, those parameters and their result are stored in the cache
- Subsequently, when the same function is invoked with the same parameter values (that is, when there is a **cache hit**), the result is retrieved from the cache, instead of being recomputed

## Result cache bypass

- In some situations, the cache is bypassed. When the cache is bypassed:
  - The function computes the result instead of retrieving it from the cache
  - The result that the function computes is not added to the cache
- Some examples of situations in which the cache is bypassed are:
  - The cache is unavailable to all sessions
  - A session is performing a DML statement on a table or view on which a result-cached function depends. Cache bypass ensures that:
    - The user of each session sees his or her own uncommitted changes
    - The PL/SQL function result cache has only committed changes that are visible to all sessions, so that uncommitted changes in one session are not visible to other sessions

## Result cache management

- The PL/SQL function result cache shares its administrative and manageability infrastructure with the Result Cache
- The database administrator manages the server result cache by specifying the `RESULT_CACHE_MAX_SIZE`, `RESULT_CACHE_MAX_RESULT` and `RESULT_CACHE_REMOTE_EXPIRATION` initialization parameters
- The `DBMS_RESULT_CACHE` package provides an interface to allow the DBA to administer that part of the shared pool that is used by the SQL result cache and the PL/SQL function result cache
- Dynamic performance views provide information to monitor the server and client result caches (`V$RESULT_CACHE_STATISTICS`, `V$RESULT_CACHE_MEMORY`, `V$RESULT_CACHE_OBJECTS`, `V$RESULT_CACHE_DEPENDENCY`)

## Functions that SQL can invoke

- To be invocable from SQL statements, a stored function (and any subprograms that it invokes) must obey the following purity rules, which are meant to control side effects:
  - When invoked from a `SELECT` statement or a parallelized `INSERT`, `UPDATE`, `DELETE`, or `MERGE` statement, the subprogram cannot modify any database tables
    - Note: if a select into statement inside a function would return a "no data found" exception, invoking this function inside SQL `SELECT` would not raise any errors but simply return `NULL` data. Invoking this function inside of a PL/SQL block would raise the "no data found" error [Discussion about this phenomenon](#)
  - When invoked from an `INSERT`, `UPDATE`, `DELETE`, or `MERGE` statement, the subprogram cannot query or modify any database tables modified by that statement
    - If a function either queries or modifies a table, and a DML statement on that table invokes the function, then `ORA-04091` (mutating-table error) occurs

- There is one exception: **ORA-04091** does not occur if a single-row **INSERT** statement that is not in a **FORALL** statement invokes the function in a **VALUES** clause
- When invoked from a **SELECT**, **INSERT**, **UPDATE**, **DELETE**, or **MERGE** statement, the subprogram cannot execute any of the following SQL statements (unless **PRAGMA AUTONOMOUS\_TRANSACTION** was specified):
  - Transaction control statements (such as **COMMIT**)
  - Session control statements (such as **SET ROLE**)
  - System control statements (such as **ALTER SYSTEM**)
  - Database definition language (DDL) statements (such as **CREATE**), which are committed automatically
- If any SQL statement in the execution part of the function violates a rule, then a runtime error occurs when that statement is parsed
- The fewer side effects a function has, the better it can be optimized in a **SELECT** statement, especially if the function is declared with the option **DETERMINISTIC** or **PARALLEL\_ENABLE**

## ✓ Work with functions (create, invoke and remove functions)

Function declaration and definition: 13.36

**CREATE FUNCTION**

**ALTER FUNCTION**

**DROP FUNCTION**

### Function declaration and definition

- A function is a subprogram that returns a value. The data type of the value is the data type of the function
- A function invocation (or call) is an expression, whose data type is that of the function
- A function declaration is also called a **function specification** or **function spec**
- Note: this topic applies to nested functions

- syntax:

```
FUNCTION function_name [(parameter_declaration [,...])] RETURN datatype --
declaration
[DETERMINISTIC | PIPELINED | PARALLEL_ENABLE | RESULT_CACHE
[relies_on_clause]]
{IS | AS} {[declare_section] body | call_spec}
```

- function declaration: Declares a function, but does not define it. The definition must appear later in the same block, subprogram, or package as the declaration
  - function heading: The function heading specifies the function name and its parameter list

- **RETURN** datatype: You cannot constrain this data type (with **NOT NULL**, for example). If datatype is a constrained subtype, then the returned value does not inherit the constraints of the subtype
- function definition: Either defines a function that was declared earlier or both declares and defines a function

## CREATE FUNCTION statement

- The **CREATE FUNCTION** statement creates or replaces a standalone function or a call specification
- A standalone function is a function (a subprogram that returns a single value) that is stored in the database
- Note: A standalone function that you create with the **CREATE FUNCTION** statement differs from a function that you declare and define in a PL/SQL block or package
- A call specification declares a Java method or a C function so that it can be invoked from PL/SQL. You can also use the SQL **CALL** statement to invoke such a method or subprogram
- Prerequisites:
  - To create or replace a standalone function in your schema, you must have the **CREATE PROCEDURE** system privilege
  - To create or replace a standalone function in another user's schema, you must have the **CREATE ANY PROCEDURE** system privilege
  - To invoke a call specification, you may need additional privileges, for example, **EXECUTE** privileges on a C library for a C call specification
  - To embed a **CREATE FUNCTION** statement inside an Oracle precompiler program, you must terminate the statement with the keyword **END-EXEC** followed by the embedded SQL statement terminator for the specific language
- syntax:

```
CREATE [OR REPLACE] [EDITIONABLE | NONEDITIONABLE] FUNCTION
plsql_function_source function_name [(parameter_declaration)]
RETURN datatype [sharing_clause] [subprogram_properties]
{IS | AS} {[declare_section] body | call_spec} ;
```

- **OR REPLACE**
  - Re-creates the function if it exists, and recompiles it
  - Users who were granted privileges on the function before it was redefined can still access the function without being regranted the privileges
  - If any function-based indexes depend on the function, then the database marks the indexes **DISABLED**
- **RETURN** datatype
  - For datatype, specify the data type of the return value of the function. The return value can have any data type supported by PL/SQL

- Note: Oracle SQL does not support invoking functions with **BOOLEAN** parameters or returns. Therefore, for SQL statements to invoke your user-defined functions, you must design them to return numbers (0 or 1) or character strings ('TRUE' or 'FALSE')
- The data type cannot specify a length, precision, or scale. The database derives the length, precision, or scale of the return value from the environment from which the function is called
- If the return type is **ANYDATASET** and you intend to use the function in the **FROM** clause of a query, then you must also specify the **PIPELINED** clause and define a describe method (**ODCITableDescribe**) as part of the implementation type of the function
- You cannot constrain this data type (with **NOT NULL**, for example)
- subprogram properties:
  - { invoker\_rights\_clause | accessible\_by\_clause | default\_collation\_clause | deterministic\_clause | parallel\_enable\_clause | result\_cache\_clause | aggregate\_clause | pipelined\_clause | sql\_macro\_clause }

## ALTER FUNCTION statement

- The **ALTER FUNCTION** statement explicitly recompiles a standalone function
- Explicit recompilation eliminates the need for implicit runtime recompilation and prevents associated runtime compilation errors and performance overhead
- Prerequisites:
  - If the function is in the **SYS** schema, you must be connected as **SYSDBA**. Otherwise, the function must be in your schema or you must have **ALTER ANY PROCEDURE** system privilege
- syntax:

```
ALTER FUNCTION function_name
{function_compile_clause | [EDITIONABLE | NONEDITIONABLE]}
```

- function compile clause

```
COMPILE [DEBUG] [compiler_parameters_clause] [REUSE SETTINGS]
```

## DROP FUNCTION statement

- The **DROP FUNCTION** statement drops a standalone function from the database
- Prerequisites
  - The function must be in your schema or you must have the **DROP ANY PROCEDURE** system privilege
- syntax:

```
DROP FUNCTION [schema.] function_name ;
```

- The database invalidates any local objects that depend on, or invoke, the dropped function
- If you subsequently reference one of these objects, then the database tries to recompile the object and returns an error if you have not re-created the dropped function
- If any statistics types are associated with the function, then the database disassociates the statistics types with the **FORCE** option and drops any user-defined statistics collected with the statistics type